

Hooman Ramezani

hooman125@gmail.com | hoomanramezani.com | [LinkedIn](#) | [GitHub](#)

EXPERIENCE

ML Specialist Solutions Architect

Jan 2025 – Present

Nebius Group

Remote

- Client-facing SA specializing in GPU performance, unblocking large-scale training and migrations to Blackwell
- Profile and optimize client workloads on B300 / GB300, raise MFU, and stand up GPU clusters (1000+ GPUs)
- Tune vLLM and SGLang and deploy optimized production inference endpoints for clients on Token Factory
- Run client proofs of concept end to end on B200 / B300 / GB300, from cluster validation to final benchmarks
- Advise AI neolabs on low-precision, parallelism and kernel-level optimizations to get the most from the latest GPUs

Machine Learning Researcher

May 2023 – Dec 2024

University Health Network

Toronto, ON

- First-author ‘LN-Transformer’ paper for lung nodule segmentation, achieving DiceC of 91.4% (SOTA)
- Built **multi-modal LLM** with unstructured text, CT slices for treatment recommendations with **84%** concordance

Machine Learning Engineer

Sept. 2022 – April 2023

Advanced Micro Devices (AMD)

Toronto, ON

- Cut **LLM inference** latency **80%** on AMD via quantization (QAT/PTQ), pruning, distillation, kernel tuning
- Benchmarked against **Triton** Inference Server (**TensorRT** backend) on k8s with dynamic batching

Machine Learning Engineer

Jan. 2022 – April 2022

DarwinAI (Acquired by Apple)

Remote

- Delivered medical vision model to Pfizer for advanced liver fibrosis diagnosis, achieving **91% sensitivity**
- Built end-to-end ML solutions for clients, with MLOps CI/CD pipelines for retraining and monitoring

PUBLICATIONS & TALKS

Distributed KV Cache for Agentic Inference on NVIDIA HGX B300 | *Nebius & WEKA Blogs*

Sep 2026

- Led the Nebius side of bringing a distributed KV cache backend to market as a product on Nebius AI Cloud
- Benchmarked 512 concurrent agentic coding sessions: **2.4x** more requests on the same GPUs, **5.4x** faster first token

MXFP8 and DeepEP for 41% Faster DeepSeek-V3 Pre-training | *PyTorch Blog*

Mar 2026

- Co-authored a joint blog post with **Meta’s PyTorch team**, presented at NVIDIA GTC 2026
- Accelerated DeepSeek-V3 (671B) pre-training **41%** on 256 B200s with MXFP8 and DeepEP, no quality loss

LN-Transformer: Lung Nodule Transformer for Sparse CT Segmentation | *CVPR*

Jun 2025

- Published a novel two-stage transformer for lung nodule segmentation featured in CVPR
- Strongest model on benchmark dataset **Dice 91.4%**, **F1 94.2%**, includes Meta SAM, DETR architectures

EDUCATION

University of Toronto

Toronto, ON

MASc in Machine Learning & Operations Research

2022 – 2024

- Thesis: Training small transformers for segmentation and detection; multimodal LLM and ViT for lung cancer

University of Waterloo

Waterloo, ON

BASc in Systems Design Engineering, GPA: 3.7 / 4.0

2017 – 2022

TECHNICAL SKILLS

Languages: Python, C/C++, CUDA, Triton, SQL, Bash

Inference: vLLM, SGLang, TensorRT-LLM, NVIDIA Dynamo, LMCache, Triton Inference Server

Training: PyTorch, TorchTitan, Megatron-LM, FSDP, DeepEP, NVSHMEM, TorchAO (MXFP8), Hugging Face

Performance & Networking: Nsight Systems, Nsight Compute, NCCL, DCGM, InfiniBand / RDMA, NVLink

Infrastructure: Kubernetes, Slurm, Docker, Terraform, Linux, Git, AWS, GCP